

Insert here your thesis' task.



**FACULTY
OF INFORMATION
TECHNOLOGY
CTU IN PRAGUE**

Master's thesis

Framework for Extraction of Wikipedia Articles Content

Oleksandr Husiev

Department of theoretical computer science
Supervisor: Ing. Milan Dojčinovski, Ph. D.

February 17, 2022

Acknowledgements

I would like to thank my supervisor Milan Dojčinovski, family and friends for their prolonged support during writing this thesis.

Declaration

I hereby declare that the presented thesis is my own work and that I have cited all sources of information in accordance with the Guideline for adhering to ethical principles when elaborating an academic final thesis.

I acknowledge that my thesis is subject to the rights and obligations stipulated by the Act No. 121/2000 Coll., the Copyright Act, as amended. In accordance with Article 46 (6) of the Act, I hereby grant a nonexclusive authorization (license) to utilize this thesis, including any and all computer programs incorporated therein or attached thereto and all corresponding documentation (hereinafter collectively referred to as the “Work”), to any and all persons that wish to utilize the Work. Such persons are entitled to use the Work in any way (including for-profit purposes) that does not detract from its value. This authorization is not limited in terms of time, location and quantity. However, all persons that makes use of the above license shall be obliged to grant a license at least in the same scope as defined above with respect to each and every work that is created (wholly or in part) based on the Work, by modifying the Work, by combining the Work with another work, by including the Work in a collection of works or by adapting the Work (including translation), and at the same time make available the source code of such work at least in a way and scope that are comparable to the way and scope in which the source code of the Work is made available.

In Prague on February 17, 2022

.....

Czech Technical University in Prague

Faculty of Information Technology

© 2022 Oleksandr Husiev. All rights reserved.

This thesis is school work as defined by Copyright Act of the Czech Republic. It has been submitted at Czech Technical University in Prague, Faculty of Information Technology. The thesis is protected by the Copyright Act and its usage without author's permission is prohibited (with exceptions defined by the Copyright Act).

Citation of this thesis

Husiev, Oleksandr. *Framework for Extraction of Wikipedia Articles Content*. Master's thesis. Czech Technical University in Prague, Faculty of Information Technology, 2022.

Abstrakt

Tato diplomová práce se zabývá extrakcí obsahu Wikipedie pro DBpedia - crowd-sourced projekt. Hlavním cílem této práce bylo vyvinout rámec pro extrakci obsahu, struktury a anotací článků z Wikipedie. Výsledkem je framework, který zpracovává velké skládky XML na Wikipedii v několika populárních jazycích s možností dynamicky přidávat nové jazyky a vytváří čistý textový výstup, odkazy a strukturu stránky ve formátu N-Triples.

Klíčová slova NIF, RDF, propojená data, web škrábání.

Abstract

This thesis describes the development process of the extraction of Wikipedia articles content for a DBpedia, a crowd-sourced community effort. The main goal of this thesis was to develop a framework for extraction of Wikipedia articles content, structure, and annotations. The result is a framework that processes large Wikipedia XML dumps in several popular languages, with the possibility to dynamically add new languages, and produces clean text output, links, and page structure in N-Triples format.

Keywords NIF, RDF, linked data, web scraping, knowledge graph.

Contents

Introduction	1
Motivation	1
Objectives	1
Challenges	2
1 Background and related works	3
1.1 The Concept of Semantic Web	3
1.2 What is Linked Data?	4
1.2.1 RDF Description	4
1.3 NLP Interchange Format	7
1.3.1 Existing Use Cases for NIF	9
1.4 Linked Open Data and DBpedia	10
1.4.1 Extracting Structured Information from Wikipedia	11
1.4.2 DBpedia Dataset	12
1.4.3 Triplestore	12
1.4.4 DBpedia Dataset Web Endpoints	13
1.5 Related works	13
1.5.1 DBpedia Information Extraction Framework	13
2 Analysis and Implementation	15
2.1 Requirements	15
2.1.1 Desired Output	16
2.1.1.1 Context	16
2.1.1.2 Links	17
2.1.1.3 Page Structure	18
2.2 Design	19
2.3 General Workflow	19
2.4 Usability considerations	20
2.4.1 REST API	20

2.4.1.1	REST API Endpoints	21
2.4.2	Command Line Interface	21
2.4.3	Command Line Input Options	22
2.5	Project Architecture	22
2.6	Implementation	24
2.6.1	Tools and libraries	24
2.6.2	Spring Framework	25
2.6.3	Java Jackson XML Library	25
2.6.4	Article Parsing	26
2.6.5	NIF Formatting	27
2.6.6	Output generation	28
2.6.7	Dynamic Language Support	28
3	Testing and Results	31
3.1	Smoke Testing	34
3.2	Unit Test coverage	34
3.2.1	JUnit Framework	35
3.3	Integration Testing	37
3.4	End-to-End Testing	39
3.4.1	SHACL Shape Validation	39
3.4.2	English language parsing	40
3.4.3	Testing other languages	42
3.4.4	Output format validation	42
3.4.5	Scale Testing	43
4	Conclusions	45
4.1	Future improvements	46
5	Acronyms	47
	Bibliography	49

List of Figures

1.1	Basic RDF Graph[1]	5
2.1	General Extraction Framework data workflow	19
2.2	Framework Main Class Diagram	24

Introduction

Motivation

Knowledge bases are growing up in importance as a Web and enterprise search engine. At the moment, knowledge bases cover only specific niches and are not useful outside of their primary purpose. As part of a broader DBpedia initiative, this thesis has an objective to structure information, store it in a machine-readable form, and provide better ways for information to be collected, organized, searched, and utilized.

A DBpedia is a knowledge base, which information is organized as an open knowledge graph. DBpedia data is served as Linked Data, which opens a new way to access the Web for applications: via browser, automated crawlers, or complex SQL-like queries. For example, current technologies do not allow to combine information about cities, criminal rates, climate, and open job postings into one search. The goal of DBpedia is to enable such queries to happen.

The creation of the Framework for Extraction of Wikipedia Articles Content is important as it will allow the DBpedia initiative to receive formatted article data from the Wikipedia on a regular basis.

Objectives

The main sight of the thesis is to extract structured content from Wikipedia articles. This content can be divided into several main parts: context, structure, and links. Context is the text itself. The structure is how the article is organized and split into sections, subsections and paragraphs, and links are either links to other Wikipedia articles or external websites. Additionally, it is essential to take care of article publication dates, clean up the non-standard articles and sections, and cover other Wikipedia languages.

1. Accept and process input data in the form of Wikipedia XML dumps

2. Extract context.
3. Extract page structure.
4. Extract links.
5. Provide outputs for context, links and page structure in the form of N-Triples.
6. Implement language extensibility.
7. Provide a user interface.

Challenges

The main problem, however, is the current size of Wikipedia. Only a full English Wikipedia dump containing only text and XML structures takes around 16 GB of space. Therefore, a thesis should research the design and implementation of not only functional but an efficient parser with both horizontal and vertical scalability in mind.

An additional challenge lies in the structure of a general wiki page. Not all pages and their components are structured in the same way. For example, some of the pages might have a line break in the middle of a citation link, that can be ambiguously interpreted as a new paragraph start. This can cause unexpected errors and exceptions during the parsing, and the goal of the project is also to minimize and gracefully handle such exceptions.

Background and related works

1.1 The Concept of Semantic Web

The Semantic Web is a Web of Data, an extension to the Web that links the related data. The collection of Semantic Web technologies (such as Resource Description Framework (RDF), Web Ontology Language (OWL), Simple Knowledge Organization System (SKOS), SPARQL Protocol and RDF Query Language) (SPARQL), etc.) provides an environment where application can query that data, draw inferences using vocabularies, etc. The goal of the Semantic Web is to develop languages to express current Internet data in a machine processable way. This is achieved by making Internet data interoperable by allowing every machine to exchange data that has unambiguous meaning defined in centralized reference resources, such as DBPedia.”Semantic Web” is a term that Tim Berners-Lee has initially coined[2] for a web of data where not simply text, but also meaning, or logical connections that machines can process. The Semantic Web is an extension of the current World Wide Web, defined by standards set by the World Wide Web Consortium (W3C).

Although Semantic Web is a term that has often been criticized as confusing, opaque, and academic, it does nonetheless capture two of the most critical aspects of these technologies:

- **Semantic:** The meaning of the data is not only explicitly represented and richly expressive, but it also “travels” along with the data itself;
- **Web:** Individual pieces of data are linked together into a network of information, just as documents are linked together on the World Wide Web.

Less familiar synonyms to the Semantic Web are the following: Linked Data Web, the Web of Data, Web 3.0, the Enterprise Information Web, or the Giant Global Graph.

1.2 What is Linked Data?

In order to make Semantic Web, or Web of Data, a reality, it is necessary to support a vast amount of data that are manageable and reachable by a variety of Semantic Web tools. Not only access to data but also relationships among data should be provided. This collection of interconnected datasets can also be referred to as Linked Data.

The term was introduced in 2006 by Tim Berners-Lee's four rules for publishing data on the Web[3], stating the following expectations for HyperText Markup Language (HTML) or RDF standards:

1. Use Uniform Resource Identifier (URI) as an identifier.
2. Use Hypertext Transfer Protocol (HTTP) URIs so that people can look up those names.
3. When someone looks up a URI, provide useful information, using the standards (RDF*, SPARQL)
4. Include links to other URIs. so that they can discover more things.

While these four rules are the basis of Linked Data and related developments, their exact implementation can be done differently and evolves over time. Specifically, the way URI is now serialized either in RDF/Extensible Markup Language (XML) or via N3(also known as Turtle, or N-triples).

1.2.1 RDF Description

The Resource Definition Framework(RDF) is a general-purpose language for representing information in the Web[4].The RDF data model is similar to classical conceptual modeling approaches, such as entity-relationship or class diagrams. It is based on the idea of making statements about the resource in expressions of the form subject–predicate–object, also known as triples. The subject denotes the resource, and the predicate denotes traits or aspects of the resource, and expresses a relationship between the subject and the object[1].

The underlying structure of any expression in RDF is a collection of triples, each consisting of a subject, a predicate and an object. A set of such triples is called an RDF graph. Each triple represents a statement of a relationship between the things denoted by the nodes that it links. Each triple has three parts - subject, predicate and object[5]. Subject and object are graph nodes, while a predicate defines a relation between them, pointing from subject to an object. On 1.1, for a triple `<http://www.w3.org/TR/rdf-syntax-grammar>` is a subject, `<http://purl.org/de/elements/1.1/title>` is a predicate and "RDF/XML Syntax Specification (Revised)" is an object.

To define what a subject is, we need to define an RDF URI Reference.

RDF URI Reference - is a Unicode string that:

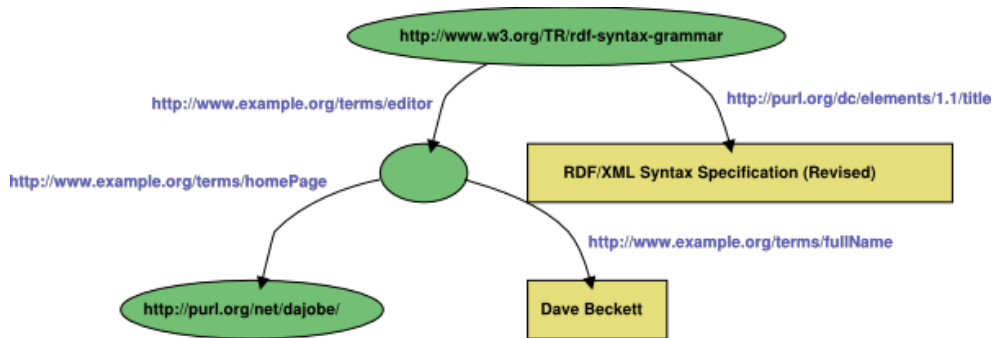


Figure 1.1: Basic RDF Graph[1]

- does not contain any control characters (`#x00 - #x1F`, `#x7F-#x9F`)
- and would produce a valid URI character sequence (per RFC2396[6], sections 2.1) representing an absolute URI with optional fragment identifier when subjected to the encoding described below.

Blank node is a special case for RDF graph nodes that do not have a URI Reference, usually represented as `"_"`. Those blank nodes represent all nodes for which graph only has partial information in a form of this node's objects.

Literal is a node that is used to identify values such as numbers, strings or dates by means of a lexical representation. Literal can be either plain or typed. For plain literals, they are represented by a text and an optional language tag. For typed literals, they are represented by a string and a type tag, for example in `<xsd:boolean, "false">` `xsd:boolean` is a type and `"false"` is a string of a type boolean.

Subject can be either an RDF URI Reference or a blank node. Predicate is an RDF URI Reference. Object can be an RDF URI Reference, Literal or a blank node.

In applications of RDF such as RDF Site Summary (RSS) resources are represented by URIs that denote and are used to link actual data on the World Wide Web[5]. RDF in general can be used not only for Internet resources, but for other knowledge ontologies too.

An RDF database, also called triplestore contains triples of interrelated statements that can be visualized with a network graph. A traditional relational database might split attributes about artworks and features about artists into separate tables. In an RDF/graph database, all these data points belong to the same interconnected graph, allowing users maximum flexibility in deciding how they wish to query it.

RDF graphs can also be represented in many text serialization formats. In a sentence "Mona Lisa was painted by Leonardo da Vinci" the subject is

1. BACKGROUND AND RELATED WORKS

"Mona Lisa", the object is "Leonardo da Vinci", and a predicate that defines a relation between the subject and the object is "paintedBy", as seen in the listing 1.1 .

Listing 1.1: Example of an RDF statement

```
<https://dbpedia.org/page/Mona_Lisa> <http://purl.org/dc/terms/
  title> "Mona Lisa" .

<https://dbpedia.org/ontology/author> <http://www.w3.org
  /1999/02/22-rdf-syntax-ns#label> "was created by" .

<https://dbpedia.org/page/Leonardo_da_Vinci> <http://xmlns.com/
  foaf/0.1/name> "Leonardo da Vinci" .
```

The particular encoding for resources or triples varies from format to format. A non-exhaustive list of popular RDF serialization formats includes:

- **Turtle** - a compact human-friendly format.
- **N3** - format that is similar to Turtle, but allows for additional features like inference, or transformation rules.
- **N-Triples** - a very simple, easy-to-parse, line-based format that is not as compact as Turtle.
- **N-Quads** - a superset of N-Triples for serializing multiple RDF graphs.
- **RDF/XML** - an XML-based syntax that was the first standard format for serializing RDF.
- **RDF/JSON** - an alternative syntax for expressing RDF triples using a simple JSON notation.
- **JSON-LD** - a JSON-based serialization, allows data to be serialized in a way that is similar to traditional JSON.

For example, it is required to write an RDF file, say `<http://example.org/smith>`, local identifiers, say `#albert`, `#brian` and `#carol`. This RDF file will look differently in XML(1.2), Turtle format(1.3). It can be observed that Turtle format is more concise than the XML one.

Listing 1.2: Example of RDF serialization in XML

```
<rdf:Description about="#albert"
  <fam:child rdf:Resource="#brian">
  <fam:child rdf:Resource="#carol">
</rdf:Description>
```

Listing 1.3: Example of RDF serialization in N3/Turtle

```
<#albert> fam:child <#brian>, <#carol>.
```

The World Wide Web (WWW) architecture now gives a global identifier "http://example.org/smith#albert" to Albert. Anyone can now use this global identifier to refer to and provide more information about Albert.

In addition to describing a link, it is essential to know when to make a link. One important pattern is a set of data which you can explore as you go link by link by fetching data. Whenever one looks up the URI for a node in the RDF graph, the server returns information about the arcs out of that node, and the arcs in. In other words, it returns any RDF statements in which the term appears as either subject or object.

Formally, call a graph *G* *browsable* if, for the URI of any node in *G*, if I look up that URI I will be returned information which describes the node, where describing a node means:

1. Returning all statements where the node is a subject or object; and
2. Describing all blank nodes attached to the node by one arc.

There are also the next limitations on such browseable data, mainly regarding data consistency across separate documents. By these definitions, statements which relate things in two different documents must be repeated. This clearly goes against the knowledge principle Don't Repeat Yourself (DRY), or in this case, not to store data in other places, as the problems with keeping the data consistent will arise eventually. A set of completely browsable data with links in both directions has to be completely consistent, and that takes coordination, especially if different authors or programs are involved.

One of the solutions to this repetition problem is to have links of a certain property in a separate document. A person's homepage doesn't list all their publications but instead puts a link to it a separate document listing them.

In conclusion, linked data is essential for linking the Semantic Web. It is quite easy to implement linked data in both new and already existing applications or websites. Various common-sense considerations determine when to make a link and when not to.

1.3 NLP Interchange Format

Natural Language Processing (NLP) is a field of science that combines linguistics, computer science and artificial intelligence concerned with the interactions between computers and human language, in particular how to program computers to process and analyze large amounts of natural language data.

The NLP Interchange Format (NIF) is an RDF/OWL-based format that aims to achieve between NLP tools, language resources and annotations[7].

NIF consists of specifications, ontologies and software that are combined under the version identifier "NIF 2.0", but are also versioned separately. The initial specification of NIF was released in November 2011.

NIF is being developed as a result and to facilitate the needs of Linked Data and related tools. NIF addresses the interoperability problem on three layers: the structural, conceptual and access layer. NIF is based on a Linked Data enabled URI scheme for identifying elements in (hyper-)texts that are described by the NIF Core Ontology (structural layer) and a selection of ontologies for describing common NLP terms and concepts (conceptual layer). NIF-aware applications will produce output adhering to the NIF Core Ontology as REST services (access layer). NIF enables the creation of heterogeneous, distributed and loosely coupled NLP applications, which use the Web as an integration platform. Another benefit is that a NIF wrapper has to be only created once for a particular tool, but enables the tool to interoperate with a potentially large number of other tools without additional adaptations. Ultimately, we envision an ecosystem of NLP tools and services to emerge using NIF for exchanging and integrating rich annotations.

NIF consists of several core components that are described below.

URI Schemes The idea behind NIF is to allow NLP tools to exchange annotations about text in RDF. Hence, the main prerequisite is that text becomes referenceable by URIs, so that they can be used as resources in RDF statements[8][9]. In NIF, there is a distinction between the document d , the text t contained in the document and possible substrings s_t of this text. We call an algorithm to systematically create identifiers for t and s_t a *URI Scheme*. The canonical URI scheme of NIF is based on RFC 5147[10], which standardizes fragment ids for the text/plain media type. According to RFC 5147, the following URI can address the first occurrence of the substring "Semantic Web" in the text (26610 characters) of the document <http://www.w3.org/DesignIssues/LinkedData.html> with the separator #: <http://www.w3.org/DesignIssues/LinkedData.html#char=717,729>.

NIF Core Ontology The NIF Core Ontology[11] provides classes and properties to describe the relations between substrings, text, documents and their URI schemes. The main class in the ontology is `nif:String`, which is the class of all words over the alphabet of Unicode characters (sometimes called Σ^*). We built NIF upon the Unicode Normalization Form C, as this follows the recommendation of the RDF standard for `rdf:Literal`. Indices are to be counted in code units. Each URI scheme is a subclass of `nif:String` and puts further restrictions over the syntax of the URIs. For example, instances of type `nif:RFC5147String` have to adhere to the NIF URI scheme based on RFC 5147. Users of NIF can create their own URI schemes by subclassing `nif:String` and providing documentation on the Web in the `rdfs:comment` field.

Framework for Extraction of Wikipedia Articles Content generates a resource of the `nif:Context` OWL class. This class is assigned to the whole string of the text (i.e. all characters). The purpose of an individual of this class is special, because the string of this individual is used to calculate the indices for all substrings. Therefore, all substrings have to have a relation `nif:referenceContext` pointing to an instance of `nif:Context`.

1.3.1 Existing Use Cases for NIF

Internationalization Tag Set The *Internationalization Tag Set (ITS)* Version 2.0 is a W3C working draft, which is in the final phase of becoming a W3C recommendation. Among other things, ITS standardizes HTML and XML attributes which can be leveraged by the localization industry (especially language service providers) to annotate HTML and XML nodes with processing information for their data value chain.

An example of three attributes in an HTML document is given here^{1.4}:

Listing 1.4: Example of Internationalization Tag Set HTML Code

```
<html>
  <body>
    <h2 translate="yes">
      Welcome to <span its-ta-ident-ref="http://dbpedia.
        org/resource/Dublin" its-within- text="yes"
        translate ="no"> Dublin </span> in <b
        translate ="no" its-within-text="yes"> Ireland
        </b>
    </h2>
  </body>
</html>
```

NIF successfully creates a bridge between ITS and RDF and a round-trip conversion was recently implemented as a proof-of-concept. Therefore, NIF can be expected to receive a wide adoption by machine translation and industrial language service providers. Additionally, the ITS Ontology provides well modeled and accepted properties, which can in turn be used to provide best practices for NLP annotations.

Ontologies of Linguistic Annotation The Ontologies of Linguistic Annotation (OLiA) provide stable identifiers for morpho-syntactical annotation tag sets, so that NLP applications can use these identifiers as an interface for interoperability. OLiA provides Annotation Models (AMs) for fine-grained identifiers of NLP tag sets. The individuals of these annotation models are then linked via `rdf:type` to coarse-grained classes from a Reference Model (RM), which provides the interface for applications. NIF provides two properties: `nif:oliaLink` links a `nif:String` to an OLiA Annotation Model. Although

a reasoner could automatically deduce the abstract type of each OLiA individual from the RM, it was a requirement that the coarse-grained types should be linked redundantly to the strings as well in case reasoning services are not available or would cause high overhead. Therefore, an OWL annotation property `nif:oliaCategory` was created as illustrated in the following example[?].

Listing 1.5: Example of Internationalization Tag Set HTML Code

```
<char=342,345> a nif:String, nif:RFC5147String;
nif:oliaLink penn:NNP;
nif:oliaCategory olia:Noun, olia:ProperNoun .
# deducable by a reasoner :
penn:NNP a olia:Noun, olia:ProperNoun .
```

1.4 Linked Open Data and DBpedia

A typical example of a large Linked Dataset is DBpedia. DBpedia and related tools are supported by the Leipzig University research group. As stated in the related article, is a community effort to extract structured information from Wikipedia and to make this information available on the Web. DBpedia allows internet users to ask sophisticated queries against datasets derived from Wikipedia and to link other datasets on the Web to Wikipedia data. This section will describe the extraction of the DBpedia datasets, and how the resulting information is published on the Web for human and machine consumption[12].

The most effective way of spurring synergistic research along these directions is to provide a rich corpus of diverse data. This would enable researchers to develop, compare, and evaluate different extraction, reasoning, and uncertainty management techniques, and to deploy operational systems on the Web. The DBpedia project has derived such a data corpus from the Wikipedia encyclopedia.

Wikipedia editions are available in over 250 languages, with the English one accounting for more than 1.95 million articles. Like many other web applications, Wikipedia has the problem that its search capabilities are limited to full-text search, which only allows very limited access to this valuable knowledge base. As has been highly publicized, Wikipedia also exhibits many of the challenging properties of collaboratively edited data: it has contradictory data, inconsistent taxonomical conventions, errors, and even spam.

The DBpedia project focuses on the task of converting Wikipedia content into structured knowledge, such that Semantic Web techniques can be employed against it — asking sophisticated queries against Wikipedia, linking it to other datasets on the Web, or creating new applications or mashups.

The DBpedia project focuses on the task of converting Wikipedia content into structured knowledge, such that Semantic Web techniques can be em-

ployed against it — asking sophisticated queries against Wikipedia, linking it to other datasets on the Web, or creating new applications or mashups. DBpedia project makes the following contributions:

- Develop an information extraction framework, which converts Wikipedia content to RDF. The basic components form a foundation upon which further research into information extraction, clustering, uncertainty management, and query processing may be conducted.
- Provide Wikipedia content as a large, multi-domain RDF dataset, which can be used in a variety of Semantic Web applications. The DBpedia dataset consists of 103 million RDF triples.
- Interlink the DBpedia dataset with other open datasets. This results in a large Web of data containing altogether around 2 billion RDF triples.
- Develop a series of interfaces and access modules, such that the dataset can be accessed via Web services and linked to other sites.

1.4.1 Extracting Structured Information from Wikipedia

Wikipedia articles contain different types of structured information, such as infobox templates, categorisation information, images, geo-coordinates, links to external Web pages and links across different language editions of Wikipedia. To process this, Wikipedia uses Mediawiki software. Due to the nature of this Wiki system, basically all editing, linking, annotating with meta-data is done inside article texts by adding special syntactic constructs. Hence, structured information can be obtained by parsing article texts for these syntactic constructs. Example of Wikipedia XML and then cleaned text that will be shown to the user can be seen in listings 1.6 and 1.7 respectively.

Listing 1.6: Raw Wikipedia XML

```
{{basic forms of government}}
'''Anarchism''' is an [[Anti-authoritarianism|anti-authoritarian
]] [[Political philosophy|political]] and [[Social
philosophy|social philosophy]]{{sfnm|1a1=McLaughlin|1y
=2007|1p=59|2a1=Flint|2y=2009|2p=27}} that rejects [[
Hierarchy|hierarchies]] deemed unjust and advocates their
replacement with [[Workers' self-management|self-managed]],
[[Self-governance|self-governed]] societies based on
voluntary, [[cooperative]] institutions. These institutions
are often described as [[Stateless society|stateless
societies]],{{sfnm|1a1=Sheehan|1y=2003|1p=85|2a1=Craig|2y
=2005|2p=14}} ...
```

Listing 1.7: Cleaned Wikipedia text

```
Anarchism is an anti-authoritarian political and social
philosophy that rejects hierarchies deemed unjust and
advocates their replacement with self-managed, self-governed
societies based on voluntary, cooperative institutions.
These institutions are often described as stateless
societies, ...
```

The XML extraction algorithm detects such Mediawiki templates and recognizes their structure using pattern matching techniques. It selects significant templates, which are then parsed and transformed to RDF triples. The algorithm uses post-processing techniques to increase the quality of the extraction. MediaWiki links are recognized and transformed to suitable URIs, common units are detected and transformed to data types. Furthermore, the algorithm can detect lists of objects, which are transformed to RDF lists.

1.4.2 DBpedia Dataset

As stated on the DBpedia's official website[13], the English version of the DBpedia dataset currently provides information about more than 4.58 million things, out of which 4.22 million are classified in a consistent ontology, including at least 1,445,000 persons, 735,000 places(including 478.000 populated places), 411,000 creative works (including 123,000 music albums, 87,000 films and 19,000 video games), 241,000 organizations (including 58,000 companies and 49,000 educational institutions), 251,000 species and 6,000 diseases.

DBpedia concepts are described by short and long abstracts in 125 languages. All these versions together describe 38.3 million things, out of which 23.8 million are localized descriptions of things that also exist in the English version of DBpedia. The full DBpedia data set features 38 million labels and abstracts in 125 different languages, 25.2 million links to images and 29.8 million links to external web pages; 80.9 million links to Wikipedia categories, and 41.2 million links to YAGO categories. DBpedia is connected with other Linked Datasets by around 50 million RDF links.

1.4.3 Triplestore

A triplestore is a software program capable of storing and indexing RDF data, in order to enable querying this data efficiently. Most triplestores support the SPARQL query language for querying RDF data. Virtuoso, Sesame, and BigOWLIM are typical examples of triplestores. DBpedia is using Virtuoso as the underlying triplestore.

1.4.4 DBpedia Dataset Web Endpoints

DBpedia website provides three access mechanisms to the DBpedia dataset: Linked Data, the SPARQL protocol, and downloadable RDF dumps. Royalty-free access to these interfaces is granted under the terms of the GNU Free Documentation License[14].

Linked Data. DBpedia resource identifiers, are set up to return RDF descriptions when accessed by Semantic Web agents, and a simple HTML view of the same information to traditional web browsers. HTTP content negotiation is used to deliver the appropriate format.

SPARQL Endpoint. Client applications can send queries over the SPARQL protocol to this endpoint at <http://dbpedia.org/sparql>. This interface is appropriate when the client application developer knows in advance exactly what information is needed. In addition to standard SPARQL, the endpoint supports several extensions of the query language that have proved useful for developing user interfaces: full text search over selected RDF predicates, and aggregate functions, notably COUNT. To protect the service from overload, limits on query cost and result size are in place. For example, a query that asks for the store's entire contents is rejected as too costly, and SELECT results are truncated at 1000 rows.

RDF Dumps. N-Triple serializations of the datasets are available for download at the DBpedia website and can be used by sites that are interested in larger parts of the dataset.

1.5 Related works

1.5.1 DBpedia Information Extraction Framework

Prior to the current project, a few projects have already been made in order to facilitate DBpedia's need for extracting information from Wikipedia and related resources, namely **DBpedia Information Extraction Framework**[15].

DBpedia Information Extraction Framework focuses on the main disadvantage of DBpedia: heavy-weight release process. Producing a DBpedia dataset release through the traditional dump-based extraction requires manual effort and – since dumps of the Wikipedia database are created on a monthly basis – DBpedia has never reflected the current state of Wikipedia. Hence, this project extended the DBpedia extraction framework to support a *live extraction*, which works on a continuous stream of updates from Wikipedia and processes that stream on the fly. More importantly, the extraction framework focuses on other parts of the Wikipedia articles. The framework has 19 extractors that process the following Wikipedia content, most important of which are list below:

1. BACKGROUND AND RELATED WORKS

- *Labels.* All Wikipedia articles have a title, which is used as an `rdfs:label` for the corresponding DBpedia resource.
- *Abstracts.* Those include a short abstract (first paragraph, represented by using `rdfs:comment`) and a long abstract (text before a table of contents, using the property `dbpedia:abstract`) from each article.
- *Interlanguage links.*
- *Images.*
- *Redirects.*
- *Disambiguation.*
- *External links.*
- *Page links.*
- *Person data.* It extracts personal information such as surname, and birth date. This information is represented in predicates such as `foaf:surname`, and `dbpedia:birthDate`.
- *Infobox*
- *Category label* Wikipedia articles are arranged in categories, and this extractor extracts the labels for those categories.

Analysis and Implementation

2.1 Requirements

Before starting to design the project, it is beneficial to specify the requirements. There several ways to layout the requirements, mainly writing down formal requirements or use cases. Because the project is oriented on delivering results to a smaller group of developers, it will be better to use formal requirements, as opposed to user-oriented use cases.

The list of functional requirements for the Framework for Extraction of Wikipedia Articles Content:

1. **Accept input data:** The framework should accept the official Wikipedia dumps in the XML format, provided by the Wikipedia. The dumps can contain an amount of information up to 20 GB of text data. The framework should be able to parse dumps in English and at least 4 other popular Wikipedia languages.
2. **Provide outputs:** The framework should print all the outputs in the N-Triples format, concatenating processed data from all articles in a single XML input file and writing the data to .nt output file.
3. **Extract context:** The framework should extract clean text from the Wikipedia page, removing or processing all the XML and Wikipedia-specific markup, including the core text but excluding infoboxes, files, images, and footers.
4. **Extract page structure:** The framework should extract a page tree, where every page section is a node, preserving the relation to the page context. The page tree should include The page tree should be printed in an output file separately from the page context.
5. **Extract links:** In addition to context, the framework should supplement the context by extracting internal Wikipedia links from the page.

These links should refer to their respective sections of the text, and include only other Wikipedia articles, excluding possible external links to other web pages. The links should also be printed in another output file, separately from the page structure and context.

6. **Implement extensibility:** The framework should be easily extensible by other developers to include new languages.
7. **Provide an interface:** The framework is required to have an intuitive interface for the user to easily leverage the framework in other works.
8. **Evaluate the results:** The framework should contain the metrics that will provide user with a feedback about every execution.

While functional requirements are the most important part of the project, there can be other, less specific requirements that go along with the functional requirements, commonly known as non-functional requirements. For this project, the non-functional requirements were next:

1. Research of Wikipedia XML Dump Structure and NIF data format.
2. Find ways to facilitate the goals of the DBpedia project.

2.1.1 Desired Output

The output that is required is defined below[16], split into context, links and page structure.

2.1.1.1 Context

Context is a resource that includes the full text of a page with all the Wikipedia formatting removed in a `nif:isString` property. This string serves as a context for its substrings. The Unicode String given in the `nif:isString` property must be used to calculate the begin and endIndex for all `nif:Strings` that have a `nif:referenceContext` property to this URI[17]. For the context resource, we need next values:

- **Predicate Language - `predLang`.** The main language of the context.
- **Context String - `isString`.** The context's text.
- **Source URL - `sourceUrl`** Link to the Wikipedia article.
- **Context's Index at the start - `beginIndex`** Starting index of the context - always starts at zero.
- **Context's Index at the end - `endIndex`** The context's length.

Listing 2.1: Example of an output for context in NIF format

```

@prefix ns0: <http://persistence.uni-leipzig.org/nlp2rdf/
  ontologies/nif-core#> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .

<http://dbpedia.org/resource/Anarchism?dbpv=2022-02&nif=context>
  ns0:predLang <http://lexvo.org/id/iso639-3/eng> ;
  ns0:sourceUrl <http://en.wikipedia.org/wiki/Anarchism> .

<http://dbpedia.org/resource/Anarchism?dbpv=2022-01&nif=context>
  ns0:isString "Anarchism is an anti-authoritarian
    political and social philosophy that rejects ..." ;
  ns0:endIndex "43346"^^xsd:nonNegativeInteger ;
  ns0:beginIndex "0"^^xsd:nonNegativeInteger ;
  a ns0:Context .

```

2.1.1.2 Links

For the links resource, we need to pick all the internal, or Wikipedia links, as well as external links that lead to other websites from Wikipedia. Note that this does not include references, only links that are part of the text. Links can be either of type *Word*, which is a single word, or *Phrase*, which includes several words that have the same link. Links should have next properties:

- **Reference Context - referenceContext.** The context resource to where this link belongs.
- **Identity Reference - taIdentRef.** Resource's identity reference in DPpedia
- **Super string - superString.** Parent paragraph of the link.
- **Anchor - anchorOf** Link's word or phrase in the text.
- **Context's Index at the start - beginIndex** Starting index of the link, as related to the context's text.
- **Context's Index at the end - endIndex** The link's end index.

Listing 2.2: Example of an output for a Word link in NIF format

```

@prefix ns0: <http://persistence.uni-leipzig.org/nlp2rdf/
  ontologies/nif-core#> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
@prefix ns1: <http://www.w3.org/2005/11/its/rdf#> .

```

```
<http://dbpedia.org/resource/Anarchism?dbpv=2022-02&nif=
word_157_170>
  a <http://persistence.uni-leipzig.org/nlp2rdf/ontologies/
nif-core#Word> ;
  ns0:referenceContext <http://dbpedia.org/resource/
Anarchism?dbpv=2022-02&nif=context> ;
  ns0:beginIndex "157"^^xsd:nonNegativeInteger ;
  ns0:endIndex "170"^^xsd:nonNegativeInteger ;
  ns0:superString <http://dbpedia.org/resource/Anarchism?
dbpv=2022-02&nif=paragraph_0_550> ;
  ns1:taIdentRef <http://dbpedia.org/resource/Self-
governance> ;
  ns0:anchorOf "self-governed" .
```

2.1.1.3 Page Structure

Page structure describes the structure of the article, with sections and paragraphs organize into a tree. Page structure has two types of resources, sections that are separated by titles in the article, and paragraphs that are part of the section.

- **Reference Context - referenceContext.** The context resource to where the page structure belongs.
- **Has Section - hasSection.** Section's children subsections.
- **Has Paragraph - hasParagraph.** Section's paragraphs.
- **First Paragraph - firstParagraph** First paragraph of the section
- **Last Paragraph - lastParagraph** Last Paragraph of the section.
- **Context's Index at the start - beginIndex** Starting index of the section, as related to the context's text.
- **Context's Index at the end - endIndex** The section's end index.

Listing 2.3: Example of an output for a Section

```
@prefix ns0: <http://persistence.uni-leipzig.org/nlp2rdf/
ontologies/nif-core#> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
```

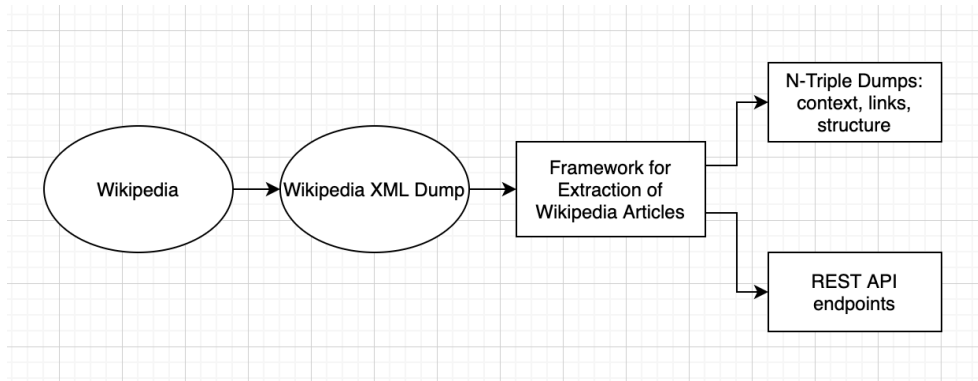



Figure 2.1: General Extraction Framework data workflow

```

<http://dbpedia.org/resource/Anarchism?dbpv=2022-02&nif=
section_0_1231>
  a <http://persistence.uni-leipzig.org/nlp2rdf/ontologies/
nif-core#Section> ;
ns0:hasSection <http://dbpedia.org/resource/Anarchism?
dbpv=2022-02&nif=section_3745_3758>, <http://dbpedia.
org/resource/Anarchism?dbpv=2022-02&nif=
section_1231_3745>, <http://dbpedia.org/resource/
Anarchism?dbpv=2022-02&nif=section_43313_43346> ;
ns0:hasParagraph <http://dbpedia.org/resource/Anarchism?
dbpv=2022-02&nif=paragraph_550_1227>, <http://dbpedia
.org/resource/Anarchism?dbpv=2022-02&nif=
paragraph_0_550> ;
ns0:referenceContext <http://dbpedia.org/resource/
Anarchism?dbpv=2022-02&nif=context> ;
ns0:lastParagraph <http://dbpedia.org/resource/Anarchism?
dbpv=2022-02&nif=paragraph_550_1227> ;
ns0:endIndex "1231"^^xsd:nonNegativeInteger ;
ns0:firstParagraph <http://dbpedia.org/resource/Anarchism
?dbpv=2022-02&nif=paragraph_0_550> ;
ns0:beginIndex "0"^^xsd:nonNegativeInteger .

```

2.2 Design

2.3 General Workflow

The general data workflow of the Framework for Extraction of Wikipedia Articles is depicted in Figure 2.1.

- Wikipedia: the main Wikipedia website is the primary sources of information.
- Wikipedia XML dump: Wikipedia runs a daily database archivation process and releases all the archived data in the form of XML dumps. These dumps can then be downloaded, unpacked and fed to the framework.
- Framework for Extraction of Wikipedia Articles: the framework processes the XML dump to get the context, structure and links and provide several options for the output.
- N-Triples Dumps: one of the outputs is to write the processed information to text files as N-Triples
- REST API endpoints: the other option is to provide the REST API endpoints for more convenient view of the output.

2.4 Usability considerations

One of the goals of the application is to make it easy to use, both by researchers and machines. In order to achieve that, application provides several interfaces. For the machines, it might be easier to connect to the application via Representational state transfer (REST) Application Programming interface (API). Humans prefer other interfaces, such as Graphic User Interface (GUI), or Command Line Interface (CLI).

2.4.1 REST API

REST is a software architecture that defines a set of constraints to be used for creating Web services. The Representational State Transfer (REST) style is an abstraction of the architectural elements within a distributed hypermedia system. REST ignores the details of component implementation and protocol syntax in order to focus on the roles of components, the constraints upon their interaction with other components, and their interpretation of significant data elements. It encompasses the fundamental constraints upon components, connectors, and data that define the basis of the Web architecture, and thus the essence of its behavior as a network-based application[18]. Web services that implement the REST architecture provide a way to communicate between different internet actors.

An API is a client-server interface which defines a contract between clients and servers.

In a RESTful Web service, requests made to a server's API Endpoint will elicit a response with a payload formatted in HTML, XML, JSON, or some other format.

The framework gives an option to use REST API endpoints over the CLI interface. While REST API is a tool used for client-server communication, DBPedia and this Framework also focus on application-based development, where no server communication is needed and all the data and tools are already downloaded and on the client side. This allows for an offline processing, as compared to constant online communication that is required if the parsing tool is only available as a server. With this in mind, we can conclude that while REST API is useful for some particular tasks, it is better to use it as a secondary interface for the Framework.

2.4.1.1 REST API Endpoints

In API terminology, communication endpoint, or simply endpoint is a unique URL address that users can access to exchange information with the server. Or in other words, APIs work using requests and responses. When an API requests information from a web application or web server, it will receive a response. The place that APIs send requests and where the resource lives, is called an endpoint[19].

The framework provides the next REST endpoints:

- **POST /articles** - Submission endpoint allows you to submit the XML dump or its part to the server. After that is done, the server will asynchronously parse the provided XML, adding the articles to the database as it goes through the submitted articles.
- **GET /articles/{title}/context** Get the context N-Triples of an article with a given title.
- **GET /articles/{title}/structure** Similarly, get the page structure of an article.
- **GET /articles/{title}/links** Get the links associated with an article with a given title.
- **GET /articles/count** Get the total count of articles in a server's database.

2.4.2 Command Line Interface

Parsing of large xml files imposes limitations on the technologies that can be used. Particularly, the size of English part of the Wikipedia xml dump has a size of 16 GB. This means that the file cannot be normally loaded into Random Access Memory (RAM), as a single modern computer will usually have from 4 to 16 GB of RAM, with Java heap utilizing a quarter of that capability by default.

Furthermore, modern internet communication is better built around frequent exchange with small packets, and imposes a limit of maximum amount of requests that can be sent in a second. For example, it will not be possible to use Wikipedia's API for this task, as the Wikipedia's server might ban all further requests. For that reason, all the processing should be done offline and not rely on the internet connection at all.

Considering the limitations described above, it was decided to use CLI as the main way to use the application.

To simplify further development process, it was decided to use an existing picocli library for simple CLI implementation, as it reduces the development time by encapsulating most of the CLI-related code, allowing the developer to control only the commands needed for CLI.

2.4.3 Command Line Input Options

The library used to create a CLI provides a good mechanisms to generate help text, from which the list of possible arguments, both mandatory and optional, can be extracted:

- **<xmlFile>** - The relative path to the XML Wiki dump.
- **-c, --clean** - The optional argument to clear output files of content before writing a new information. Useful option for testing the framework.
- **-h, --help** - Show this help message and exit. Help Screen is an important part of the CLI that lists all other commands.
- **-l, --language=<language>** - Provide the language of the XML dump that is being parsed. Default language is English.
- **-o, --output=<outputPath>** - The NIF files output folder.
- **-V, --version** - Print framework version information. This is important as different versions of the Framework can provide known bugs. Printing the version allows user to track those bugs.

2.5 Project Architecture

Since the related works mentioned in 1.5.1 are based on Java and other Java Virtual Machine (JVM) based technologies such as Scala, it is best to build the project on those technologies in order to leverage the existing knowledge and reuse already developed libraries where possible. It is, however, worth noting that a large amount of dependencies used will introduce more complexity into the system. When the dependency's behavior that is not controlled by the framework is changed, the whole application workflow may be affected

until the fix is implemented. Therefore, using only necessary dependencies is important.

Following the Java application standards, the codebase was split into packages, with **org.dbpedia** as the main package prefix, common for all DBpedia-related projects. A package in Java is used to group related classes, and is similar to a folder or a directory. We use packages to avoid name conflicts, and to write a better maintainable code.

- **application** - package for main Java classes. The application is developed in such a way that it has several main methods, and depending on the configuration, only one of those will be used.
- **cli** - package responsible for generating CLI.
- **configuration** - package that contains necessary Spring configuration, further described in ??.
- **exception** - package responsible for custom exception handling. While Java has its own Exception handling system, for a better handling it is recommended to extend the existing interfaces and catch custom exceptions instead.
- **extractor** - main package that contains all the code responsible for the extraction, parsing and structuring the information.
- **splitter** - support package that helps to split the xml dump into separate pages.

After designing the package structure, it is important to identify the general class diagram outlay. This outlay is presented in Figure 2.2. The classes and their purpose are broken down below:

- **ExtractionApplicationCLI** - main executable class that contains Java's main method.
- **XmlInput** - a class that processes the CLI input parameters, by using the picocli library[20].
- **LanguageIdentifierBean** - a singleton class that defines the language of an XML dump that is processed, set to English by default. A singleton classes can only have one instance and usually contains project-wide settings. Singleton mechanic is handled by the Spring framework, described in section ??.
- **XmlDumpService** - a service class that is a wrapper for all XML-processing operations.

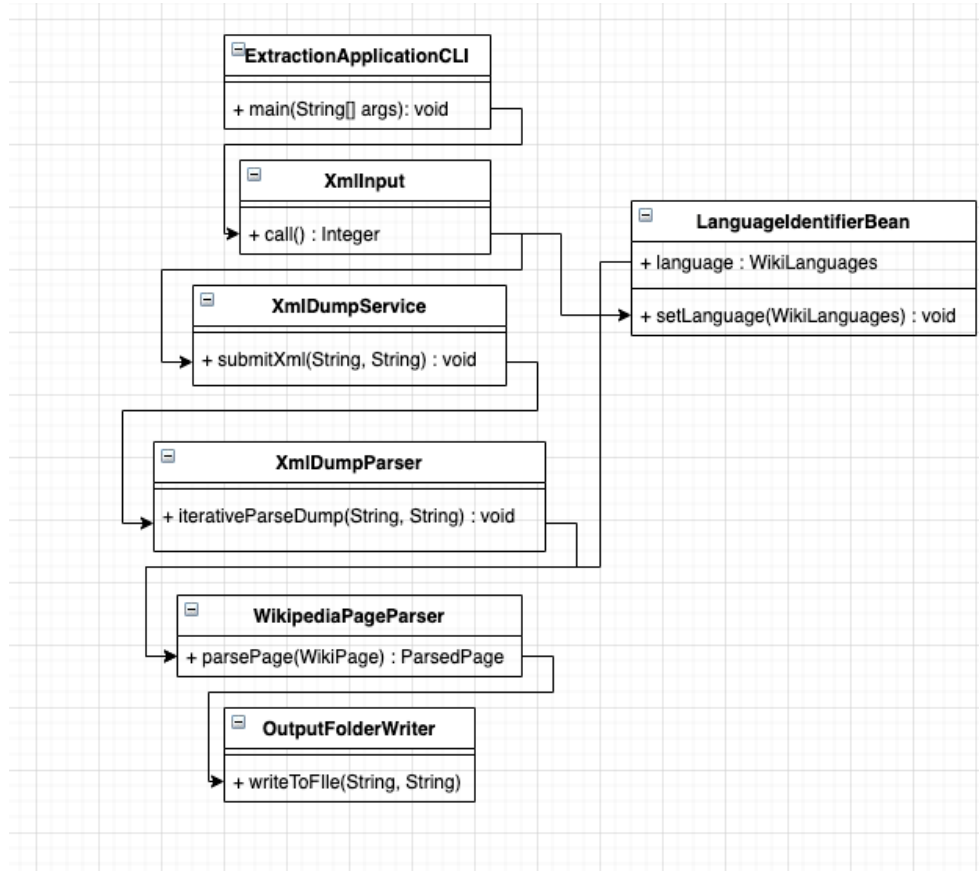


Figure 2.2: Framework Main Class Diagram

- **XmlDumpParser** - a specific class that processes XML dump and breaks it up into pages.
- **WikipediaPageParser** - a parser class that focuses on processing a single page.
- **OutputFolderWriter** - a writer class that facilitates the output of a current project.

2.6 Implementation

2.6.1 Tools and libraries

The chosen language for framework development is Java, as it is used in DBpedia Extraction Framework, described in section 1.5.1. It is, however, important to notice that plain Java is not sufficient to achieve the goal of

implementing modern Web Framework. For that reason, many other supplementary tools, libraries and other technologies were developed. This current ecosystem will be described below.

The project was developed using an IntelliJ IDEA[21], a modern Integrated Development Environment (IDE) that is maintained by a Czech company JetBrains. The company is now also developing a JVM-based language called Kotlin, that will be able to overcome some of Java's own shortcomings, such as a lack of support for functional programming paradigm. For the current project, however, it was decided to eliminate unnecessary dependencies and use Java as a main development language. IntelliJ IDEA was used for development and testing of the framework.

2.6.2 Spring Framework

The Spring Framework is a Java platform that provides comprehensive infrastructure support for developing Java applications. Spring handles the infrastructure so you can focus on your application. Spring enables you to build applications from "plain old Java objects" (Plain Old Java Object (POJO)) and to apply enterprise services non-invasively. This capability applies to the Java Standard Edition programming model and to full and partial Java Enterprise Edition[23].

There is a several versions of Spring Framework, namely the original Spring and a newer version called Spring Boot. The difference between those two is that the original Spring leverages the use of XML configurations, while Spring Boot instead uses Java Configuration classes. For the development, Spring Boot is used.

For this project, mostly prototype and singleton Bean scopes were used, with singleton being a default one.

2.6.3 Java Jackson XML Library

The Jackson project is a collection of data processing tools for the Java language and the JVM platform. It supports a wide range of data formats such as Comma-separated values (CSV), Java Properties, XML, and Yet ANother Markup Language (YAML) through extension components that support the specific language.

The Jackson XML component is meant for reading and writing XML data by emulating how Jakarta XML Binding (JAXB) works, although not conclusively.

In this project, Jackson library is used to serialize Java objects into XML and deserialize them back into Java objects, in order to eventually produce text output.

XmlMapper Class XmlMapper is the main class from Jackson 2.x that helps the developers in serialization. This mapper is available in jackson-dataformat-xml jar, that can be easily added to the project using Apache Maven - project's dependency management[24].

An example of deserialization can be seen in Listing 2.4. Here, a Mediawiki class is a POJO class, or in other words a simple mapping of an XML scheme to a Java class hierarchy, where every subcomponent of a schema is mapped to a Java subclass.

Listing 2.4: Example of an XML Deserialization

```
private XmlMapper xmlMapper = new XmlMapper();

private Mediawiki deserializeXml(String dump) throws
    IOException {
    return xmlMapper.readValue(dump, Mediawiki.class);
}
```

2.6.4 Article Parsing

The most important part of the Framework is a parsing component. While the code utilizes Jackson XML Library 2.6.3 to deserialize XML, the article text is processed in a WikipediaPageParser class. First, we need to clean up the text from many tags used by Wikipedia[25]. Tags and patterns listed above utilize regular expressions to match and remove patterns in the text. The exact list of tags removed is described in the next list:

- **Gallery.** Both tags(<gallery/>) and their contents are removed.
- **Unit Conversion.** For this tag, the Framework extracts 2nd and 3rd parts and separates them by a space, i.e. for ”{{convert|2|km|mi}}” output is ”2 km”
- **Emphasis.** Removes single-quote tags from text(”’ or ”) while keeping its contents intact.
- **No Table of Contents.** Remove technical ”__NOTOC__” tags that hide table of contents.
- **Indentation.** Replaces excessive line breaks with simple n.
- **Math formula.** Remove all math formulas(<math>) and their contents.
- **IPA.** Removes all International Phonetic Alphabet(IPA) tags.

Regular expression is done using Java's Pattern class. In the example of a regular expression 2.5 that is used to remove gallery, the regular expression checks for a pattern `<gallery>...</gallery>`, ignoring cases and including line breaks as part of the `.*` regular expression, passing `Pattern.DOTALL` parameter.

Listing 2.5: Example of a regular expression

```
private static final Pattern GALLERY = Pattern.compile("&lt;gallery&gt;.*?&lt;/gallery&gt;[\\n]?",
    Pattern.CASE_INSENSITIVE | Pattern.DOTALL);
```

2.6.5 NIF Formatting

After the article is parsed, the `ParsedPage` object is passed to the `NifFormatter` generator class methods to format output lines. This class utilizes Apache Jena modelling to create RDF models, resources and properties. In the code listing 2.6, we can find the code that generates RDF resources for DBpedia and context resource, and then proceeds to create context's type property (`http://persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-core#Context`) and its starting index (`"0"<http://www.w3.org/2001/XMLSchema#nonNegativeInteger>`).

Listing 2.6: Generating Apache Jena context model

```
Resource dbPediaResource = jenaModel.createResource(dbpediaUrl);
Resource contextResource = jenaModel.createResource(
    PERSISTENCE_ONTOLOGY_LINK + "#" + LinkType.CONTEXT.
    getCapitalizedTypeLabel());

// Context NIF type
Property rdfSyntaxProperty = jenaModel.createProperty(
    RDF_SYNTAX_TYPE);
dbPediaResource.addProperty(rdfSyntaxProperty, contextResource);
// Context beginIndex
Property beginIndexProperty = jenaModel.createProperty(
    PERSISTENCE_ONTOLOGY_LINK, "#" + BEGIN_INDEX);
dbPediaResource.addProperty(beginIndexProperty, jenaModel.
    createTypedLiteral(beginIndex, XSDDatatype.
    XSDnonNegativeInteger));
```

`NifFormatter` converts parse page objects into Apache Jena models. For each article, three different models are produced: context model, page structure model and links model. Many of the strings required to generate the model are hardcoded, such as links(DBpedia URL, ontology) or property names.

Context and Link model creation is straightforward, but page structure is more complicated. Page structure generation is done using a recursive function. Main method is `generatePageStructureEntry`, and this method calls a recursive function `generateNodeEntry` for the article's root section. Then a recursive call processes every paragraph of the section, adding them as section resource properties with property name "hasParagraph". First and last paragraph have to be added as separate properties, so first and last paragraphs are added twice with "firstParagraph" and "lastParagraph" properties as well as again with "hasParagraph". After processing all section's paragraph, every subsection of a section is then processed by recursive call. Subsections then are constructed with the same algorithm, until all the article's sections are processed.

2.6.6 Output generation

`OutputFolderWriter` is a simple class that utilizes `java.io` to produce output in the desired folder. The path to the folder is passed via `-output` CLI variable, and `OutputFolderWriter` creates three files in the folder: `nif_context.nt` for article's context resources, `nif_links.nt` for article's links and `nif_structure.nt` for article's section structure. If the output folder is not created yet, the framework will create it. When the framework processes several articles, `OutputFolderWriter` iteratively appends output of every article to the files.

2.6.7 Dynamic Language Support

The XML-structure of the Wikipedia article does not differ much from language to language. There are only a few points to be aware of: Footer headings and categories. In English, those would be "See also", "References", "Further reading", "External Links", and "Related pages", and Categories simply have a heading "Category". Those are the parts that have to be removed from the articles.

One of the project's requirements was to implement an easy extensibility mechanism to add new languages, mention in the Requirements Section 2.1. This was achieved by adding an abstract class `LanguageFooterRemover`. This class has some general functions to parse parts of the text that might be unique to different languages. Before the framework starts, it processes the configuration file `language_list.xml` that is stored in the configuration folder. The examples of this file's contents can be seen in Listing 2.7. Such template can easily be reused to extend the number of supported languages if needed.

Listing 2.7: Example of an language configuration file

```
<languageContainer>
  <language>
    <langName>ENGLISH</langName>
```

```
        <categoryName>Category</categoryName>
        <footer>See also</footer>
        <footer>References</footer>
        <footer>Further reading</footer>
        <footer>External Links</footer>
        <footer>Related pages</footer>
    </language>
    <language>
    <langName>POLISH</langName>
        <categoryName>Kategoria</categoryName>
        <footer>Przypisy</footer>
        <footer>Uwagi</footer>
    </language>
</languageContainer>
```

Testing and Results

The main machine that was used for testing, unless specified otherwise, had next specifications:

- **Processor:** 2,7 Hz Dual-Core Intel Core i5;
- **Memory:** 8 GB 1867 MHz DDR3;
- **Graphics Card:** Intel Iris Graphics 6100 1536 MB;
- **Operating System:** macOS Catalina, Version 10.15.6.

Project benchmarks **Benchmark Testing** measures a repeatable set of quantifiable results that serves as a point of reference against which products/services can be compared. The purpose of benchmark testing results is to compare the present and future software releases with their respective benchmarks.

A benchmark must be repeatable. For instance, with every iteration of load a test, if the response times varies too much, system performance be benchmarked. Response time needs to be stable amongst different load conditions.

A benchmark must be quantifiable. For example, the user experience cannot be quantified in numbers, but time a user spends on a webpage due to good UI can be quantified.

Next benchmark tests were repeatedly used to test the framework. They are located in the project's folder documents:

- 1 Page test, English language. Expected to be completed quickly and successfully.
- 257 pages, English language. Execution time is expected to be less than a minute.

3. TESTING AND RESULTS

- 6738 pages, English language. Execution time is expected to be around 5 minutes.
- 6649 pages, Polish language. Execution time is expected to be around 5 minutes.
- 15118 pages, English language. Execution time is around 25 minutes.

Benchmarking can be used not only for performance testing, but also for testing the quality of output. While testing this project, the results were constantly compared to an ideal output, that has been provided by the DBpedia project[26]. Dbpedia website contains ideal output format on the resources page at <https://www.dbpedia.org/resources/nif/>. The output on the page is shown for the Antropology page in Turtle format, but it can be used to compare output properties for context, links and page structure.

Three different output components were compared separately. On the listings below, the ideal result can be seen, for the context output in Listing 3.1 - the output containing properties for the context. For links in Listing 3.2, the example contains output for one of the phrases for "Austroasiatic languages" page. For page structure in Listing 3.3, the output shows an example of a Section and its reference to the parent context resource.

Listing 3.1: Exemplary result for NIF Context

```
<http://dbpedia.org/resource/Anarchism?dbpv=2020-02&nif=context>
  <http://persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-
core#isString> "Anarchism is a radical political ... \n*
Textbooks from Wikibooks \n* Data from Wikidata \n* Anarchy
Archives. Anarchy Archives is an online research center on
the history and theory of anarchism" .
```

Listing 3.2: Exemplary result for NIF Links

```
<http://dbpedia.org/resource/Austroasiatic_languages?dbpv
=2016-04&nif=phrase_94_109> <http://www.w3.org/1999/02/22-
rdf-syntax-ns#type> <http://persistence.uni-leipzig.org/
nlp2rdf/ontologies/nif-core#Phrase> .
<http://dbpedia.org/resource/Austroasiatic_languages?dbpv
=2016-04&nif=phrase_94_109> <http://persistence.uni-leipzig.
org/nlp2rdf/ontologies/nif-core#referenceContext> <http://
dbpedia.org/resource/Austroasiatic_languages?dbpv=2016-04&
nif=context> .
<http://dbpedia.org/resource/Austroasiatic_languages?dbpv
=2016-04&nif=phrase_94_109> <http://persistence.uni-leipzig.
org/nlp2rdf/ontologies/nif-core#beginIndex> "94"^^<http://
www.w3.org/2001/XMLSchema#nonNegativeInteger> .
```

```

<http://dbpedia.org/resource/Austroasiatic_languages?dbpv
=2016-04&nif=phrase_94_109> <http://persistence.uni-leipzig.
org/nlp2rdf/ontologies/nif-core#endIndex> "109"^^<http://www
.w3.org/2001/XMLSchema#nonNegativeInteger> .
<http://dbpedia.org/resource/Austroasiatic_languages?dbpv
=2016-04&nif=phrase_94_109> <http://persistence.uni-leipzig.
org/nlp2rdf/ontologies/nif-core#superString>

```

Listing 3.3: Exemplary result for NIF Page Structure

```

<http://dbpedia.org/resource/Ada?dbpv=2016-04&nif=section_0_17>
<http://www.w3.org/1999/02/22-rdf-syntax-ns#type> <http://
persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-core#
Section> .
<http://dbpedia.org/resource/Ada?dbpv=2016-04&nif=section_0_17>
<http://persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-
core#beginIndex> "0"^^<http://www.w3.org/2001/XMLSchema#
nonNegativeInteger> .
<http://dbpedia.org/resource/Ada?dbpv=2016-04&nif=section_0_17>
<http://persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-
core#endIndex> "17"^^<http://www.w3.org/2001/XMLSchema#
nonNegativeInteger> .
<http://dbpedia.org/resource/Ada?dbpv=2016-04&nif=section_0_17>
<http://persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-
core#referenceContext> <http://dbpedia.org/resource/Ada?dbpv
=2016-04&nif=context> .
...

```

The output format should be in N-triples[27], better described in Section 1.2.1.

3.1 Smoke Testing

Smoke test is a test or a test suite that covers the main functionality of a component or system to determine whether it works properly before planned testing begins[28].

To list the information, those are the advantages of an early and continuous smoke testing:

- It exposes any integration issues, that is if there any problems in communication between different tools used in the project.
- It uncovers problems early,
- It provides some level of confidence that changes to the software have not adversely affected major areas (the areas covered by smoke testing)

3. TESTING AND RESULTS

This kind of testing was performed during every stage of the implementation after each new part of functionality has been added to the framework. It includes CLI testing, API testing, functionality testing, and output verification. During development and testing the framework was running locally on a Mac-based machine.

A big amount of bugs and errors was revealed and subsequently fixed during these tests. For example, there were many errors related to the parsing of Wikipedia articles, and lots of inconsistencies happening during the page structure translation into the page.

Furthermore, scaling the framework has caused a lot of problems. The size of an English Wikipedia dump is about 16 GB of data, and parsing it takes a lot of time. During those tests, it was discovered that such amount of data can not be handled by the server, and therefore an adequate API for production purposes can not be easily provided.

For the smoke testing, the next tests were conducted:

- Single-article XML dump in English language.
- Two-page XML dump in English language.
- Other languages support testing, conducted for a single page from a German segment of Wikipedia.

3.2 Unit Test coverage

Unit testing is a software testing method by which individual units of source code - sets of one or more computer program modules together with associated control data, usage procedures, and operating procedures - are tested to determine whether they are fit for use. The goal of unit testing is to isolate each part of the program and show that the individual parts are correct. A unit test provides a strict, written contract that the piece of code must satisfy. As a result, it affords several benefits[29].

Unit testing finds problems early in the development cycle. This includes both bugs in the programmer's implementation and flaws or missing parts of the specification for the unit. The process of writing a thorough set of tests forces the author to think through inputs, outputs, and error conditions, and thus more crisply define the unit's desired behavior. The cost of finding a bug before coding begins or when the code is first written is considerably lower than the cost of detecting, identifying, and correcting the bug later. Bugs in released code may also cause costly problems for the end-users of the software. Code can be impossible or difficult to unit test if poorly written, thus unit testing can force developers to structure functions and objects in better ways[30][31][32].

An important indicator for the quality of tests is their code coverage. Code coverage is measured by what percentage of the code is tested by unit tests. Running an analysis shows that 73% of the current Framework lines, 53% of its methods, or 52% of its classes are covered by tests. Both unit and integration tests are included into this statistic, as the analysis tool does not differentiate those tests.

3.2.1 JUnit Framework

For the framework implementation, a JUnit library was used. **JUnit** is a popular Testing Framework used to implement unit tests in Java[33].

For a given project, unit tests were used to test the execution of a `WikipediaPageParser` class, that is used to parse separate pages, as well as its supplementary classes, such as a `DumpSplitService`. You can see the examples of a unit test used in the project in the Listing 3.4:

Listing 3.4: JUnit Paragraph Parsing Unit Test Class

```
@Log4j
public class WikipediaPageParserTest {

    private static WikipediaPageParser pageParser;
    private static WikiPage wikiPage;
    private static XmlTransformer contextLanguageTransformer;

    @BeforeAll
    public static void beforeAll() throws IOException {
        pageParser = new WikipediaPageParser(new
            ContextLanguageTransformer());

        URL textUrl = Resources.getResource("page_test.txt");
        wikiPage = new WikiPage("Anarchism",
            Resources.toString(textUrl, StandardCharsets.UTF_8));
    }

    @Test
    public void parseParagraphsTest() throws IOException,
        ParsingException {
        Subdivision root = pageParser.buildPageStructure(
            wikiPage);
        // check that the paragraphs are parsed
        assertTrue(root.getParagraphs().size() > 1);
        // check that the page has a meaningful structure
```

3. TESTING AND RESULTS

```
        assertTrue(root.getChildren().size() > 1);
    }
    ...
}
```

Those tests are checking for specific methods to return some basically expected values, as well as test possible corner cases that were previously discovered as bugs, but were already fixed later. Such tests can be classified as regression tests. For example, Listing 3.5 is a test that checks whether `<ref>` tags are removed from the text by the `WikiTagsRemover` class. Those tags were initially causing issues with parsing, and a unit test has been added to check whether they are properly removed. If the `WikiTagsRemover` method is broken during development, the test will immediately indicate that method's output is now broken (or that a "regression" has happened). Those tests are run on every project recompilation.

Listing 3.5: Regression Testing

```
@Test
public void removeRefTags() {
    String refText = "&lt;ref&gt;{{cite OED|anarchism}}&lt;/ref&gt; and the <ref>word</ref> &lt;ref&gt;{{cite OED|anarchism}}&lt;/ref&gt;";
    refText = StringEscapeUtils.unescapeHtml4(
        StringEscapeUtils.unescapeHtml4(refText));

    Pattern HTML_TAGS = Pattern.compile("<[^\>]+>");
    refText = wikiTagsRemover.removeHtmlTags(refText);
    log.info(refText);
    Assert.assertFalse(HTML_TAGS.matcher(refText).find());
}
```

3.3 Integration Testing

Integration testing is the phase in software testing in which individual software modules are combined and tested as a group. Integration testing is conducted to evaluate the compliance of a system or component with specified functional requirements[34]. It occurs after unit testing and before system testing. Integration testing takes as its input modules that have been unit tested, groups them in larger aggregates, applies tests defined in an integration test plan to those aggregates, and delivers as its output the integrated system ready for system testing[35].

Integration testing in the project tests that all aspects of the Framework work as intended, from parsing to output, and different cases with varying sizes

and languages are automatically tested whenever the project is recompiled. To speed up the development, it is possible to skip those tests by compiling the project with `-DskipTests` argument, but this is not recommended. Integration tests compare Framework's output for prepared articles that are stored in the test folder with expected values that are manually verified to be true.

Integration tests are placed in the `OutputValidationTests` class. Every test is verifying any specific information:

- **test*ArticleGeneral** - tests that the total amount of nodes/lines in the context, structure and links files is correct;
- **test*ArticleContext** - tests context output file to verify that every node and property is in the output;
- **test*ArticleLinks** - tests if the links output file is formed correctly. Since there is a lot of links in every article, only selected links are tested;
- **test*ArticleStructure** - tests if the page structure output file is formed correctly. Similarly to links, this test verifies that a selected part of sections are correctly formed in the output.

Apache Jena is used to parse the output and create a model. The model is then verified by the integration test, as can be seen on Listing 3.6:

Listing 3.6: Integration Testing - English Article Context

```
@Test
public void testEnglishArticleContext() throws IOException {
    loadArticle("xml_test_page.xml", "ENGLISH");

    File outputContext = Paths.get(folder.getRoot().
        getAbsolutePath(), OutputFolderWriter.
        CONTEXT_FILENAME).toFile();

    jenaModel.read(new FileInputStream(outputContext), null,
        NTRIPLES);

    String testSubject = String.format("http://dbpedia.org/
        resource/Anarchism?dbpv=%s&nif=context"
        , getCurrentDateString());
    String testPredicate = "http://www.w3.org/1999/02/22-rdf-
        syntax-ns#type";
    String testObject = getPersistenceOntologyUrl("Context");

    // test if type is Context
    assertStatementContains(testSubject, testPredicate,
        testObject);
}
```

```
.....

// test predLang
assertStatementContains(testSubject, testPredicate,
    testObjectLiteral);

testPredicate = getPersistenceOntologyUrl("sourceUrl");
testObjectLiteral = createLiteral("http://en.wikipedia.
    org/wiki/Anarchism");

//test sourceUrl
assertStatementContains(testSubject, testPredicate,
    testObjectLiteral);

jenaModel.removeAll();
}
```

Integration tests also verify that other languages are properly processed. Currently integration tests cover English and German languages. For a German article, integration test on Listing 3.7 verifies that language input parameter is processed correctly and the context output specifies the correct language:

Listing 3.7: Integration Testing - German Article Context

```
@Test
public void testGermanArticleContext() throws IOException{
    loadArticle("xml_test_deutsch_one_page.xml", "GERMAN");
    ....

    testPredicate = getPersistenceOntologyUrl("predLang");
    Literal testObjectLiteral = createLiteral("http://lexvo.
        org/id/iso639-3/deu");

    // test predLang
    assertStatementContains(testSubject, testPredicate,
        testObjectLiteral);
}
```

3.4 End-to-End Testing

End-to-end testing is a Software testing methodology to test an application flow from start to end. The purpose of End-to-end testing is to simulate the

real user scenario and validate the system under test and its components for integration and data integrity. What it means for this project is that we will need to do a test from downloading an XML dump from Wikipedia to receiving the processed results. Here is the general outlay of a testing process:

1. Download the latest Wikipedia dump[36]. They are released at least monthly and usually twice a month. Different languages are listed after the metawiki dumps.
2. Unzip the file, in Unix-like operating systems usually done in `bzip2 -d wikidatawiki-*-pages-meta-history1.xml-p1p224.bz2`

Optional Because parsing of the whole XML dump is time-consuming, I have used a way to reduce the amount of articles that is processed. To do it, it is possible to execute this command `head -n 100000 enwiki-20191101-pages-articles-multistream1.xml > short_test_100k.xml`, and then properly close off the XML by removing the last article and close the XML brackets.

3. Build the framework and pass the input parameters to a file. I have created a script `parse_xml_dump.sh` that will build the framework if necessary and run the executable with the XML dump path as the parameter.

3.4.1 SHACL Shape Validation

To ensure that RDF resources produced are not malformed, Shapes Constraint Language (SHACL) forms were utilized. SHACL shapes allow to define certain classes for RDF resources and then to validate any RDF resources against those shapes, making it possible to run the tests on any output. In the Listing 3.8 a definition for a Context class is shown. SHACL definitions are defined in a Turtle format, which is another format from NIF format, produced by the framework.

Listing 3.8: SHACL Context Class definition

```
ontology:Context
rdf:type rdfs:Class ;
rdfs:label "Context" ;
.

ontology:isString
rdf:type rdf:Property ;
rdfs:domain ontology:Context ;
rdfs:label "Context string" ;
rdfs:range xsd:string ;
```

```
.  
...  
ontology:ContextShape  
a sh:NodeShape ;  
sh:targetClass ontology:Context ;  
sh:property ontology:isStringShape ;  
sh:property ontology:sourceUrlShape ;  
sh:property ontology:predLangShape ;  
sh:property ontology:beginIndexShape ;  
sh:property ontology:endIndexShape ;  
.  
  
ontology:isStringShape  
a sh:isStringShape ;  
sh:path ontology:isString ;  
sh:minCount 1 ;  
sh:maxCount 1 ;  
sh:datatype xsd:string ;  
sh:severity sh:Violation  
.
```

To validate output against SHACL shapes, various tools can be used, notably an open-source validator to test shapes and data: <https://shacl.org/playground/>. Application's output can be converted to a Turtle format using EasyRDF Converter from <https://www.easyrdf.org/converter/>.

Current SHACL shapes validate whether properties are present and that their number does not exceed the required number. For more complex validation, integration tests are used.

3.4.2 English language parsing

Initially the testing was done on a single English-language article. This testing uncovered many problems with the parsing model, such as the need to update the recursive function to build the page structure. Most importantly, the first testing helped to understand the vast number of Wikipedia XML components. Of those, I had to drop the parts that were already covered by the previous frameworks mentioned in Section 1.5.1, such as infoboxes, images, files, categories and others, and instead focus only on the text. I also dropped the citations mentioned in the article's footer, and instead focused on the text itself.

In the Listing 3.9, an example of the Framework's output is shown. Every run shows the total amount of pages parsed, as well as what percentage of articles was not parsed due to Framework's exceptions. In this case 96.74% of articles were parsed and added to the output. If there is a processing

error, it is indicated in the output by providing the reason for an error and part of the paragraph where the error happened, for example "Error parsing page Crankshaft: Broken xml component - closing brace not found for {{ in paragraph {{Main|Science and t..."

Listing 3.9: Example of an English processing command

```
macs-MacBook-Air:wiki-realtime-extractor mac$ ./parse_xml_dump.
  sh documents/large_documents/short_test_1000k.xml --language
    =ENGLISH --clean -o=output
....
22-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page: Cemetery
  H culture
2022-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page:
  Corrado Gini
2022-01-31 00:08:36 ERROR XmlDumpParser:121 - Error parsing page
  Crankshaft: Broken xml component - closing brace not found
  for "{" in paragraph "{{Main|Science and t..."
2022-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page:
  Central nervous system
2022-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page: Caste
2022-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page:
  Creation
2022-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page: Coral
  66
2022-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page:
  Rhyming slang
2022-01-31 00:08:36 INFO XmlDumpParser:112 - Parsed page:
  Canchim
2022-01-31 00:08:36 INFO XmlDumpParser:135 - Total pages parsed:
  3378. Success rate: 96.74%. Seconds passed: 198
```

3.4.3 Testing other languages

Additionally to English language, I have added other popular Wikipedia languages. According to the latest Wikipedia statistics, those are the main languages of Wikipedia[37]:

1. **English:** 2,567,509 articles, 22.5% of the total number of articles;
2. **German:** 808,044 articles, 7.1%;
3. **French:** 709,312 articles, 6.2%;
4. **Polish:** 539,688 articles, 4.7%;

3. TESTING AND RESULTS

5. Japanese: 523,629 articles, 4.6%.

Dynamic Language Support is described in 2.6.7. To verify that the Framework works in other languages, I have downloaded English, German and Polish XML dumps and run the framework against them, as seen on Listing 3.10.

Listing 3.10: Example of a Polish processing command

```
macs-MacBook-Air:wiki-realtime-extractor mac$ ./parse_xml_dump.  
sh documents/xml_test_polish.xml --language=POLISH --clean -  
o=output  
2022-01-31 00:00:35 INFO XmlDumpParser:112 - Parsed page: AWK  
2022-01-31 00:00:36 INFO XmlDumpParser:112 - Parsed page:  
Alergologia  
2022-01-31 00:00:37 INFO XmlDumpParser:112 - Parsed page:  
Aksjomat  
2022-01-31 00:00:37 INFO XmlDumpParser:112 - Parsed page:  
Arytmetyka  
....  
2022-01-31 00:00:40 INFO XmlDumpParser:112 - Parsed page:  
Alabama  
2022-01-31 00:00:40 INFO XmlDumpParser:135 - Total pages parsed:  
76. Success rate: 100.00%. Seconds passed: 5
```

3.4.4 Output format validation

There are several utilities that can be used to validate the output, most notable Apache Jena and rapper[38]. For the output validation, I have picked rapper, as it is more lightweight and has the ability to count or parse the provided N-triples. For example, the command `rapper --input ntriples --output rdxml --show-graphs nif_links.nt` will parse the links and transform it into RDF/XML format. If this transformation will not throw any exceptions, this will mean that the output is a valid set of N-triples. Rapper tool can also be used to convert output into a different format, such as RDF.

3.4.5 Scale Testing

For the scale testing, I have implemented logging and a simple parsing success metric, as some of the articles have a syntax that may deviate from the standard or the framework's programmed expectations. For example, some might contain links that are broken with line separators, or differently encoded XML components.

To further measure the time that the application will take to parse the code, I have added the execution time metric. To make sure it is still possible to execute, I added a `-dry-run` CLI option that cleans up the output files to

remove output after every execution. Not doing that results in generating text that takes a lot of disk space. Unfortunately, output clean up slows down execution by around 3 times, so this option is unusable for a general performance testing.

I have run several tests over the English Wiki dump with the next results. The table shows results with appending and keeping text in output files, that is faster because the file output connection is not closed after every article parsing.

# of pages parsed	Total execution time	ms per article	Success rate
257	26	110	100%
257	46	178	100%
2087	109	50	88.5%
6738	237	35	87.9%
6649	414	63	96.35%
15118	1433	94	96.47%

Conclusions

The result of this thesis is a Java Framework that allows users to parse and retrieve the Wikipedia XML dump and achieves most of the original objectives, in some places with a room for improvement:

1. **Accept and process input data in the form of Wikipedia XML dumps.** The Wikipedia XML Dump parsing was achieved, and the process to do so best described in Section 3.4. The statistics show that the parsing success rate averages on 88% over the large amounts of articles, meaning that around 12% of articles will contain some kind of component that will not be parsable by the framework and will be skipped. These systemic errors can be avoided by further investigation of Wikipedia's XML Format.
2. **Extract context.** Context is extracted and stored in the form of N-Triples. Some of the contexts might still contain traces of the original XML code. This can be later fixed by improving the XML removal code.
3. **Extract page structure.** Page structure is extracted and recursively built in the form of N-Triples.
4. **Extract links.** Links are extracted, URLs that link them to the page structure are created.
5. **Provide outputs for context, links and page structure in the form of N-Triples.** Output is printed.
6. **Implement language extensibility.** Language extensibility mechanism is implemented, new languages can be added in the form of an XML that is parsed into POJO when the application is starting, better described in Section 2.6.7. Framework is tested in several languages - English, German and Polish.

7. **Provide a user interface.** User interface is provided in two different forms and is described in Section 2.4.

4.1 Future improvements

There is a number of future improvements that can be done to this project:

- Code quality improvements. There is currently some unused code, style warnings and duplicate fragments that can be updated to clean up the code.
- Article parsing success rate. By analyzing the logs, more articles can be successfully parsed by adding more rules to regular expression editing of the context.
- Add parsing of Wikipedia footers. Only the main article text is currently parsed by the Framework, the footers, specifically citations, might contain useful information.
- Add horizontal scalability. Execution on the main Wikipedia dump files takes a lot of time, with 25 minutes for each XML part and 27 parts available for English, the whole processing can take around 10 to 11 hours. The work, however, can be easily parallelized by splitting the input to be processed by different machines.

Acronyms

API Application Programming interface.

CLI Command Line Interface.

CSV Comma-separated values.

DRY Don't Repeat Yourself.

FOAF Friend Of A Friend.

FR Functional Requirement.

GUI Graphic User Interface.

HTML HyperText Markup Language.

HTTP Hypertext Transfer Protocol.

IDE Integrated Development Environment.

J2EE Java 2 Platform Enterprise Edition.

JAXB Jakarta XML Binding.

JVM Java Virtual Machine.

NIF NLP Interchange Format.

NLP Natural Language Processing.

OWL Web Ontology Language.

POJO Plain Old Java Object.

RAM Random Access Memory.

RDF Resource Description Framework.

REST Representational state transfer.

RSS RDF Site Summary.

SKOS Simple Knowledge Organization System.

SPARQL SPARQL Protocol and RDF Query Language).

URI Uniform Resource Identifier.

W3C World Wide Web Consortium.

WWW World Wide Web.

XML Extensible Markup Language.

YAML Yet ANother Markup Language.

Bibliography

- [1] Beckett, D. RDF/XML Syntax Specification. In *W3C Recommendation - RDF/XML Syntax Specification*, February 2004.
- [2] Berners-Lee, T. Q&A with Tim Berners-Lee [online]. April 2007. Available from: <https://www.bloomberg.com/news/articles/2007-04-09/q-and-a-with-tim-berners-leebusinessweek-business-news-stock-market-and-financial-advice>
- [3] Berners-Lee, T. Linked Data [online]. July 2006. Available from: <https://www.w3.org/DesignIssues/LinkedData.html>
- [4] Roussey, C.; Pinet, F.; et al. *An Introduction to Ontologies and Ontology Engineering*. London: Springer London, 2011, ISBN 978-0-85729-724-2, pp. 9–38, doi:10.1007/978-0-85729-724-2_2. Available from: https://doi.org/10.1007/978-0-85729-724-2_2
- [5] Klyne, G. RDF Concepts and Abstract Syntax. In *W3C Recommendation - RDF Concepts and Abstract Syntax*, February 2004.
- [6] Berners-Lee, T. Uniform Resource Identifiers (URI): Generic Syntax. RFC 2396, RFC Editor, August 1998. Available from: <https://www.ietf.org/rfc/rfc2396.txt>
- [7] Sebastian Hellmann, S. A., Jens Lehmann; Brümmer, M. Integrating NLP using Linked Data. In *12th International Semantic Web Conference*, Sydney, Australia, October 2013.
- [8] Hellmann, S.; Lehmann, J.; et al. Linked-Data Aware URI Schemes for Referencing Text Fragments. 10 2012, ISBN 978-3-642-33875-5, pp. 2–4, doi:10.1007/978-3-642-16438-5_10.
- [9] Hellmann, S.; Lehmann, J.; et al. Integrating NLP using Linked Data. 10 2013, ISBN 978-3-642-38708-1, doi:10.1007/978-3-642-41338-4_7.

- [10] E. Wilde, M. D. URI Fragment Identifiers for the text/plain Media Type.
- [11] NIF 2.0 Core Ontology [online]. 2020. Available from: <https://persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-core/nif-core.html>
- [12] Auer, S.; Bizer, C.; et al. DBpedia: A Nucleus for a Web of Open Data. 01 2007, ISBN 978-3-540-76297-3, pp. 722–735, doi:10.1007/978-3-540-76298-0_52.
- [13] DBpedia Facts & Figures [online]. September 2020. Available from: <https://dbpediawww.informatik.uni-leipzig.de/about/facts-figures>
- [14] Morsey, M.; Lehmann, J.; et al. DBpedia and the live extraction of structured data from Wikipedia. *Program: electronic library and information systems*, volume 46, 04 2012: pp. 157–181, doi:10.1108/00330331211221828.
- [15] Hellmann, S.; Stadler, C.; et al. DBpedia Live Extraction. 11 2009, ISBN 978-3-642-05150-0, pp. 1209–1223, doi:10.1007/978-3-642-05151-7_33.
- [16] DBpedia Language Resources(NIF). January 2022. Available from: <https://www.dbpedia.org/resources/nif/>
- [17] NIF Core Ontology. January 2022. Available from: <https://persistence.uni-leipzig.org/nlp2rdf/ontologies/nif-core/nif-core.html>
- [18] Fielding, R. T. Architectural Styles and the Design of Network-based Software Architectures.
- [19] Kapadnis, J. REST: Good Practices for API Design[online]. March 2018. Available from: <https://medium.com/hashmapinc/rest-good-practices-for-api-design-881439796dc9>
- [20] Picocli - command line interface library [online]. 2020. Available from: <https://picocli.info/>
- [21] IntelliJ IDEA - Capable and Ergonomic IDE [online]. 2020. Available from: <https://www.jetbrains.com/idea/>
- [22] JetBrains Java development ecosystem research [online]. 2020. Available from: <https://www.jetbrains.com/lp/devecosystem-2020>
- [23] Spring Framework Reference Documentation. Available from: <https://docs.spring.io/spring-framework/docs/4.3.12.RELEASE/spring-framework-reference/htmlsingle/>

- [24] Apache Maven Project [online]. 2020. Available from: <https://maven.apache.org/pom.html>
- [25] Wikipedia Help:Template. Available from: <https://en.wikipedia.org/wiki/Help:Template>
- [26] DBpedia Project Link [online]. 2020. Available from: <https://wiki.dbpedia.org/>
- [27] N-Triples - a line-based syntax for an RDF graph [online]. 2020. Available from: <https://www.w3.org/TR/n-triples/>
- [28] Smoke Testing. Available from: <https://softwaretestingfundamentals.com/smoke-testing/>
- [29] Kolawa, A.; Huizinga, D. *Automated Defect Prevention: Best Practices in Software Management*. IEEE Computer Society Press. p. 75. -0-470-04212-0, 2007, ISBN 978.
- [30] Boehm, B. W.; Papaccio, P. N. Understanding and Controlling Software Costs. *IEEE Transactions on Software Engineering*, volume 14, no. 10, May 2016: pp. 1462–1477, doi:10.1109/32.6191.
- [31] Test Early And Often. Available from: [https://docs.microsoft.com/en-us/previous-versions/visualstudio/visual-studio-2012/ee330950\(v=vs.110\)?redirectedfrom=MSDN](https://docs.microsoft.com/en-us/previous-versions/visualstudio/visual-studio-2012/ee330950(v=vs.110)?redirectedfrom=MSDN)
- [32] Prove It Works: Using the Unit Test Framework for Software Testing and Validation. Available from: <https://www.ni.com/en-us/innovations/white-papers/09/prove-it-works--using-the-unit-test-framework-for-software-testi.html>
- [33] JUnit 5 User Guide. Available from: <https://junit.org/junit5/docs/current/user-guide/>
- [34] Standard - Systems and software engineering. *ISO/IEC/IEEE*, volume 15, December 2010: pp. 1–418.
- [35] Ould, M. A.; (ed), C. U. Testing in Software Development, BCS (1986), p71. *Accessed*, volume 31, October 2014.
- [36] Wikimedia Downloads [online]. 2020. Available from: <https://www.w3.org/TR/n-triples/>
- [37] Top Ten Wikipedias [online]. 2020. Available from: https://meta.wikimedia.org/wiki/Top_Ten_Wikipedias
- [38] Raptor RDF Syntax Library - Raptor RDF parser utility [online]. 2014. Available from: <http://librdf.org/raptor/rapper.html>